Split	Method	AUC@20	AUC@50	AUC@100	Final
Train-6	MAGICIAN	0.326	0.592	0.761	**0.945**
	Ours (V3-T0 + D1)	**0.438**	**0.662**	**0.795**	0.940
	Δ (Ours − MAGICIAN) [95% CI]	+0.111 [+0.028, +0.195]	+0.071 [+0.012, +0.128]	+0.034 [+0.004, +0.064]	-0.005 [-0.010, -0.001]
Dev-5	MAGICIAN	0.317	0.543	0.695	0.883
	Ours (V3-T0 + D1)	**0.439**	**0.645**	**0.760**	**0.889**
	Δ (Ours − MAGICIAN) [95% CI]	+0.122 [+0.073, +0.167]	+0.102 [+0.018, +0.195]	+0.064 [-0.018, +0.162]	+0.006 [-0.077, +0.072]
Test-3	MAGICIAN	0.252	0.452	0.638	0.867
	Ours (V3-T0 + D1)	**0.355**	**0.557**	**0.707**	**0.888**
	Δ (Ours − MAGICIAN) [95% CI]	+0.103 [+0.044, +0.162]	+0.105 [+0.058, +0.157]	+0.070 [+0.045, +0.088]	+0.021 [-0.005, +0.064]
Dev+Test-8	MAGICIAN	0.293	0.509	0.674	0.877
	Ours (V3-T0 + D1)	**0.407**	**0.612**	**0.740**	**0.888**
	Δ (Ours − MAGICIAN) [95% CI]	+0.115 [+0.078, +0.150]	+0.103 [+0.048, +0.166]	+0.066 [+0.012, +0.130]	+0.011 [-0.041, +0.058]

Same episodes as Table 1. Splits are those of our model (T0 was trained on Train-6 only; Dev-5 / Test-3 were never used for training or checkpoint selection).

They are 'T0-held-out' scenes, not scenes unseen by every method: MAGICIAN's (MACARONS) weights were trained on 8 of the 14 scenes (split_manifest.csv), and

Test-3 results of earlier planner variants had been inspected, so Test-3 is not a blind test. Δ: scene-paired difference, scene bootstrap 95% CI (few scenes: wide).