Split	Method	AUC ↑			Coverage ↑		
		@20	@50	@100	@20	@50	Final
Train (6)	MAGICIAN	0.326	0.592	0.761	0.579	0.877	**0.945**
	Ours	**0.438**	**0.662**	**0.795**	**0.662**	**0.897**	0.940
Dev (5)	MAGICIAN	0.317	0.543	0.695	0.532	0.784	0.883
	Ours	**0.439**	**0.645**	**0.760**	**0.669**	**0.845**	**0.889**
Test (3)	MAGICIAN	0.252	0.452	0.638	0.460	0.703	0.867
	Ours	**0.355**	**0.557**	**0.707**	**0.574**	**0.782**	**0.888**
Dev+Test (8)	MAGICIAN	0.293	0.509	0.674	0.505	0.753	0.877
	Ours	**0.407**	**0.612**	**0.740**	**0.634**	**0.821**	**0.888**

Scene-equal means (5 official starts per scene); Final = coverage after 100 steps; bold: better of the two.

Our network was trained on Train scenes only (checkpoint selected on held-out Train trajectories); the planner variant was chosen on Dev; Test was used for neither. MAGICIAN's weights were trained on 8 of the 14 scenes (all Train scenes, 2 Dev scenes).